

Strategic Alignment and Gap Analysis of the GIMAN Research Plan with the PPMI High-Impact Research Roadmap

Executive Summary

Overview

This report presents a comprehensive assessment of the Graph-Informed Multimodal Attention Network (GIMAN) research plan, evaluating its alignment with the strategic intent of the Parkinson's Progression Markers Initiative (PPMI) High-Impact Research Roadmap. The analysis finds that the GIMAN plan proposes a technically innovative and methodologically sophisticated framework for prognostic modeling in Parkinson's disease (PD). Its core architectural design, which leverages Graph Neural Networks (GNNs) and advanced attention mechanisms for deep multimodal fusion, is exceptionally congruent with the roadmap's strategic goals for methodological advancement. However, the report identifies a critical strategic misalignment in the plan's application. By focusing exclusively on a *de novo* PD cohort and defining static prognostic endpoints, the GIMAN plan overlooks several key research frontiers, most notably the roadmap's highest-priority area: the prediction of disease onset in the prodromal cohort.

Key Findings

The assessment reveals a mixed landscape of alignment between the GIMAN plan and the five strategic frontiers defined by the PPMI roadmap:

- **Strong Alignment:** The plan demonstrates exceptional alignment with **Frontier 2 (Deep**

Multi-Modal Fusion) and **Frontier 5 (Explainable AI)**. The proposed "gradual fusion" architecture is a state-of-the-art implementation of the principles advocated in the roadmap, and its multi-scale interpretability framework represents a robust strategy for biomarker discovery.

- **Partial Alignment:** The plan shows partial alignment with **Frontier 1 (Dynamic Progression Modeling)** and **Frontier 3 (Deconstructing Heterogeneity)**. While GIMAN's proposed prognostic endpoints are clinically relevant, they are static summaries of progression and lack the dynamic, time-to-event nature called for by the roadmap. Similarly, its approach to subtyping is implicit and post-hoc, falling short of the roadmap's vision for modeling heterogeneity as a primary, continuous signature.
- **Critical Gap:** The plan has a complete and critical gap concerning **Frontier 4 (The Prodromal Frontier)**. The GIMAN plan's exclusive focus on a cohort of already-diagnosed PD patients represents a fundamental divergence from the PPMI's most significant strategic pivot toward understanding and predicting disease onset in at-risk individuals.

Top-Level Recommendations

To address these gaps and elevate the GIMAN project's strategic impact, this report puts forth several key recommendations. The most crucial of these is a major reframing of the project to include a parallel GIMAN model specifically designed for the PPMI prodromal cohort, with the primary objective of predicting phenoconversion. Further recommendations include incorporating a time-to-event (survival analysis) framework to predict clinically meaningful disability milestones and integrating the alpha-synuclein Seed Amplification Assay (SAA) biomarker as a biological ground truth to validate the model's learned representations.

Concluding Statement

The GIMAN research plan details a technically brilliant and powerful methodology. However, to transition from a strong methodological contribution into a landmark study that fully realizes the PPMI's vision for transformative Parkinson's research, it requires significant strategic redirection. By implementing the recommendations outlined in this report, the GIMAN project can align its formidable technical capabilities with the most pressing scientific and clinical questions in the field, maximizing its potential for high-impact discovery and accelerating the path toward precision medicine in Parkinson's disease.

Introduction: Juxtaposing Strategic Vision with Proposed Execution

The PPMI Strategic Imperative

The Parkinson's Progression Markers Initiative (PPMI) has evolved from a data repository into a dynamic research ecosystem with a clear strategic vision: to identify and validate biological markers that track PD progression, with the ultimate goal of accelerating the development of disease-modifying therapies.¹ A critical review of the research landscape catalyzed by PPMI reveals a significant disconnect between this core mission and the predominant focus of the machine learning community. A 2023 review found that the vast majority of studies focused on diagnostic classification (distinguishing PD from healthy controls), a task of limited clinical utility for the already-diagnosed PPMI cohort.¹ The review explicitly concludes that "much of what makes the PPMI data set unique (multi-modal and longitudinal observations) remains underutilized in most machine learning studies".¹ The PPMI Strategic Roadmap was developed as a direct response to these limitations, outlining five key research "frontiers" that represent a clear mandate for future high-impact research. These frontiers call for a decisive shift away from static, unimodal, diagnostic tasks and toward dynamic, multimodal, prognostic models that can deconstruct the heterogeneity of PD, predict its onset, and provide interpretable insights.¹

The GIMAN Methodological Innovation

The GIMAN research plan is founded on a central hypothesis that directly confronts a core limitation of existing models: the assumption that each patient is an independent and identically distributed (i.i.d.) data point.³ The plan proposes a paradigm shift from purely patient-centric prediction to a population-informed framework, postulating that "Modeling the latent structure of the patient cohort as a graph and leveraging this structure...will lead to more robust, generalizable, and interpretable prognostic models".³ The core innovation lies in the application of Graph Neural Networks (GNNs) to a patient similarity graph, a method designed to explicitly model the rich, latent relationships within a patient cohort.³ This

approach allows the model to learn not only from a single patient's data but also from the data of their most similar peers, mimicking the reasoning of a clinical expert.³ By proposing to overcome the limiting i.i.d. assumption that constrains virtually all current approaches, the GIMAN plan positions itself as a methodologically sophisticated proposal designed to address the complex data interdependencies inherent in the PPMI dataset.³ This report will now assess how effectively this proposed execution aligns with the strategic vision laid out in the PPMI roadmap.

Frontier 1 Alignment Analysis: Dynamic Modeling of Disease Progression

Roadmap Mandate: Modeling the Full Trajectory

The first frontier of the PPMI roadmap calls for a fundamental pivot from static, cross-sectional analyses to dynamic models that capture the evolution of PD over time.¹ The mandate is to move beyond predicting a single outcome at a distant future point and instead model the

entire trajectory of a patient's clinical and biological journey.¹ This approach requires focusing on the

rate of change (the mathematical derivative) of biomarkers and clinical scores, rather than their absolute values, as progression is by definition change over time.¹ A central component of this frontier is the prediction of clinically meaningful disability milestones. A recent analysis of the PPMI cohort identified 25 such milestones across six domains, including events like "requiring a walking aid" or "MoCA score falling below 21".¹ The roadmap explicitly advocates for the use of survival analysis, or time-to-event modeling, to predict the timing of these events, suggesting techniques like the Cox Proportional Hazards model or more flexible machine learning approaches such as Random Survival Forests.¹ This establishes a clear benchmark: true dynamic modeling involves predicting

when a clinically significant event will occur for an individual, providing a continuous risk profile over time.

GIMAN's Proposed Endpoints: A Static Snapshot of Progression

The GIMAN research plan defines two primary prognostic tasks to be addressed by the model³.

1. **Motor Progression (Regression Task):** The model will be trained to predict the rate of motor decline, operationalized as the slope of the MDS-UPDRS Part III score calculated via linear regression for each patient over a 36-month period. The output is a single continuous variable representing the predicted annual change in UPDRS-III points.³
2. **Cognitive Decline (Classification Task):** The model will be trained to predict the risk of significant cognitive decline, framed as a binary classification task: predicting whether a patient will convert to Mild Cognitive Impairment (MCI) or dementia within a fixed 36-month follow-up window.³

Assessment of Alignment: Partial

The GIMAN plan demonstrates partial alignment with this frontier. The proposed endpoints are clinically relevant and represent a necessary step beyond simple diagnostic classification, aligning with the roadmap's general call to focus on progression.¹ However, these endpoints fall significantly short of the roadmap's vision for true dynamic modeling. A calculated slope is a linear, static summary of what is often a complex, non-linear progression curve, discarding all information about the shape and fluctuations of the trajectory. Similarly, a binary classification for a fixed 36-month window answers

if an event might happen within that period but fails to capture the crucial temporal dimension of *when* it is most likely to occur. This static approach limits the clinical actionability of the predictions, as it cannot provide the continuous, time-dependent risk profile advocated by the roadmap.

The Disconnect Between Method and Endpoint

A deeper analysis reveals a notable disconnect between the GIMAN plan's methodological capabilities and its chosen endpoints. The architectural blueprint specifies the use of sophisticated sequential encoders, such as a hybrid 3D CNN followed by a Gated Recurrent Unit (GRU), to model the spatiotemporal evolution of neuroimaging data and a separate GRU

to map the trajectory of clinical scores.³ These recurrent neural network components are explicitly designed to process sequences of data over time and learn the underlying dynamics of progression.³

This demonstrates that the proposed GIMAN architecture is technically capable of handling and learning from rich time-series data. The model is designed to internally "see" the dynamic movie of a patient's progression. However, the final output of the model is constrained to a single, static value—either a slope or a binary probability. This creates a functional inconsistency: the model's sophisticated internal machinery for processing temporal dynamics is ultimately used only to produce a static summary. The limitation is not in the model's architectural design, which is powerful and forward-thinking, but in the definition of its prognostic endpoints. The plan effectively builds a movie projector but tasks it only with displaying a single photograph.

Strategic Recommendation: Embrace Time-to-Event Analysis

To bridge this gap and fully align with Frontier 1, the GIMAN plan should be revised to incorporate a time-to-event (survival analysis) framework. This would leverage the existing architectural strengths of the model to address a more dynamic and clinically actionable question.

A new or revised aim (e.g., Aim 2b) should be introduced to implement this framework. Instead of terminating in a simple multi-layer perceptron (MLP) for regression or classification, the GIMAN architecture can be adapted to output the parameters of a survival distribution or a time-dependent hazard function. This can be readily achieved by integrating established deep survival models, such as **DeepSurv**, a deep learning extension of the Cox model, or **DeepHit**, which directly learns the survival distribution without the proportional hazards assumption.⁷

By adopting this approach, the revised GIMAN model could predict an individualized risk curve for reaching any of the 25 clinically meaningful milestones identified in the PPMI cohort, such as the 5-year risk of developing freezing of gait or the probability of requiring a walking aid over time.¹ This change would transform the model's output from a static prediction into a dynamic, personalized forecast, dramatically increasing its clinical utility for patient counseling, care planning, and trial enrichment, thereby achieving full alignment with the strategic intent of this frontier.

Frontier 2 Alignment Analysis: Deep Multi-Modal

Fusion

Roadmap Mandate: Holistic Integration and Biological Grounding

Frontier 2 of the roadmap emphasizes that PD is a multi-system disorder that cannot be understood through a single data modality.¹ It calls for the development of advanced fusion architectures that move beyond simple "early fusion" or concatenation to intelligently model the complex inter-dependencies between different biological scales—from genomics and proteomics to neuroimaging and clinical symptoms.¹ The roadmap specifically highlights methodologies such as joint embedding, intermediate fusion, and, most powerfully, cross-modal attention mechanisms that allow for dynamic, context-dependent integration of information.¹

Critically, this frontier also introduces the alpha-synuclein Seed Amplification Assay (SAA) as a "revolutionary development" in the PD field.¹ The SAA provides a direct, objective biochemical signature of the core molecular pathology of PD with remarkably high accuracy.¹⁴ The roadmap thus proposes a high-impact project: leveraging SAA data as a biological "ground truth" label to train multimodal models, with the goal of developing non-invasive proxies for underlying synuclein pathology.¹

GIMAN's Architecture: A State-of-the-Art "Gradual Fusion" Pipeline

The GIMAN architecture represents an exemplary implementation of the advanced fusion principles advocated by the roadmap. The plan details a sophisticated, multi-stage "gradual fusion" pipeline that intelligently integrates information at progressively more abstract levels³:

1. **Intra-Modal Temporal Fusion:** The plan first employs modality-specific encoders, such as the CNN-GRU for imaging and a separate GRU for clinical data, to learn the temporal dynamics *within* each data stream.³ This corresponds to the concept of intermediate fusion.
2. **Inter-Patient Graph-based Fusion:** The high-level feature vectors from the encoders are then used to initialize nodes in the patient similarity graph. A Graph Attention Network (GAT) propagates information across the graph, fusing information *across similar patients* to refine and regularize each individual's representation.³

3. **Inter-Modal Attention-based Fusion:** Finally, a cross-modal attention module is applied to the graph-informed patient embedding. This module learns to dynamically assign importance weights to each of the feature streams (imaging, genomic, clinical), fusing information *across modalities* in a context-dependent manner before making a final prediction.³

Assessment of Alignment: Strong (Architecturally), with a Critical Biomarker Omission

From an architectural standpoint, the GIMAN plan is exceptionally well-aligned with Frontier 2. Its hierarchical, gradual fusion pipeline is a state-of-the-art design that directly embodies the advanced methodologies called for in the roadmap. It moves far beyond simple concatenation and provides a clear framework for modeling complex inter-dependencies.

However, the plan has a critical omission: it makes no mention of the SAA biomarker.³ The PPMI dataset now includes SAA results, including binary positive/negative status and kinetic parameters, and the roadmap explicitly flags this data as a transformative opportunity for grounding computational models in the core biology of the disease.¹ By failing to incorporate this key biomarker, the GIMAN plan misses a major opportunity to validate its learned features and enhance its scientific impact.

Missed Opportunity for Biological Validation and Non-Invasive Proxy Development

The roadmap explicitly proposes using SAA status as a "ground truth" label for the core molecular pathology of PD—the presence of misfolded alpha-synuclein.¹ This provides a powerful alternative to training models on clinical endpoints, which are downstream consequences of the pathology and can be ambiguous or subject to measurement error. The GIMAN framework, designed to integrate non-invasive data (MRI) and minimally invasive data (genetics, blood-based proteomics, clinical scores), is perfectly positioned to tackle the high-impact project proposed in the roadmap: building a model that predicts a patient's SAA status from more accessible data types.¹

Success in this task would yield two major benefits. First, it would biologically ground the model's learned representations. By training the model to identify the signature of synucleinopathy, the features it learns are more likely to be directly related to the underlying

disease process. Second, it would lead to the development of a powerful, scalable, non-invasive proxy biomarker for the core pathology of PD. Given that CSF collection for SAA is an invasive procedure that limits its widespread use, a validated model that can predict SAA status with high accuracy from MRI, genetics, and clinical data would be a tool of immense clinical value, particularly for screening participants for clinical trials of alpha-synuclein-targeting therapies.¹ The GIMAN plan possesses the ideal technical machinery for this task but fails to recognize and seize the opportunity.

Strategic Recommendation: Incorporate SAA as a Primary or Secondary Endpoint

To rectify this omission and significantly enhance the project's impact, the GIMAN plan should be revised to incorporate the SAA biomarker. A secondary objective should be added to Aim 2: **to train and validate the GIMAN model to predict a patient's binary SAA status using the non-CSF data modalities as input.**

This revision directly addresses a key strategic project outlined in the roadmap and leverages a high-value, newly available dataset within PPMI.¹ The successful development of a non-invasive proxy for synucleinopathy could be a more significant scientific and clinical contribution than a marginal improvement in predicting UPDRS scores. Furthermore, this new objective would synergize powerfully with the proposed interpretability framework in Aim 4. The Explainable AI (XAI) tools could be applied to this SAA-prediction model to deconstruct its decisions and identify the specific imaging patterns, genetic markers, and clinical signs that serve as the strongest non-invasive proxies for the core molecular pathology of PD. This would transform the project from a pure prediction task into a powerful engine for scientific discovery, fully aligning with the deepest intent of Frontier 2.

Frontier 3 Alignment Analysis: Deconstructing Disease Heterogeneity

Roadmap Mandate: From Coarse Subtypes to Continuous Signatures

Frontier 3 confronts the consensus that PD is not a monolithic entity but a deeply heterogeneous disorder.¹ The roadmap calls for research to move beyond coarse, clinically defined subtypes (e.g., "tremor-dominant") and toward the discovery of robust, biologically-grounded patient subgroups that emerge directly from the data.¹ High-impact studies have already used unsupervised clustering on PPMI data to identify reproducible subtypes with distinct progression rates (e.g., "slow," "moderate," and "fast" progressors).¹

However, the roadmap identifies the true research frontier as moving *beyond discrete clusters* to a more sophisticated paradigm. It argues that patient heterogeneity likely exists on a continuum and proposes representing it not as a categorical label but as a position in a **continuous, low-dimensional "latent space"**.¹ The proposed methodology involves applying advanced dimensionality reduction techniques, such as

Variational Autoencoders (VAEs), to fused, multi-modal patient data to learn a personalized "progression signature" for each individual.¹ This continuous representation can capture the subtle, multi-faceted nature of a patient's disease profile far more effectively than a single subtype label.

GIMAN's Approach: Implicit Subtyping via Graph Structure

The GIMAN research plan addresses the challenge of heterogeneity implicitly through its core graph-based methodology. The construction of the patient similarity graph in Aim 1 is explicitly designed to "capture underlying patient subtypes and disease structures".³ The plan then proposes to uncover these structures in a post-hoc fashion as part of its interpretability framework in Aim 4. Specifically, it suggests using GNNExplainer, a graph attribution technique, to identify "influential patient archetypes" and "patient subgraphs" that are important for the model's predictions.³ In this framework, subtypes are not a primary learning objective but rather an emergent property of the learned graph to be explored after the model is trained for its primary prognostic tasks.

Assessment of Alignment: Partial

GIMAN's graph-based approach is a powerful and innovative method for exploring heterogeneity, and its goal of identifying data-driven patient communities is well-aligned with the initial part of Frontier 3. However, the alignment is only partial. The plan treats subtyping as a secondary, post-hoc interpretation task rather than a primary, unsupervised learning

objective. More significantly, it does not propose any method for generating the continuous, personalized "progression signatures" that the roadmap identifies as the cutting edge of heterogeneity research.¹ The GIMAN plan aims to find clusters but does not provide a framework for modeling the continuous landscape in which those clusters exist.

The Untapped Potential of the Graph Embedding

The GIMAN architecture, as designed, is uniquely positioned to create the very data needed for continuous heterogeneity modeling. After the Graph Attention Network (GAT) layer propagates information across the patient similarity graph, it produces a rich, graph-informed, multimodal feature vector—an embedding—for each patient.³ This embedding is the final representation that is fed to the prediction head.

By its very design, this embedding is a powerful, low-dimensional representation that synthesizes a patient's individual biomarker trajectory, their full multimodal data profile, and their relationship to their most similar peers in the cohort. It is a data-driven signature of that patient's disease state within the context of the population. The roadmap calls for using techniques like VAEs to learn a continuous latent space that represents heterogeneity.¹ The patient embeddings generated by GIMAN are an ideal, highly-processed input for just such a model. The GIMAN plan already computes the necessary data but stops short of explicitly modeling its structure to reveal the continuous landscape of the disease. The graph provides a powerful, regularized embedding that could serve as a superior foundation for learning a robust and interpretable latent space compared to using raw data alone.

Strategic Recommendation: Add an Unsupervised Aim for Latent Space Modeling

To fully align with Frontier 3 and unlock the full potential of its own methodology, the GIMAN plan should be revised to include a new aim dedicated to the unsupervised deconstruction of heterogeneity. A new Aim (e.g., Aim 1b or a new Aim 5) should be proposed with the following objective: **to model the continuous landscape of PD heterogeneity by applying a Variational Autoencoder (VAE) to the final patient embeddings generated by the trained GIMAN model.**

This new aim would proceed in two stages. First, the primary GIMAN model would be trained for its prognostic tasks as planned. Second, the final patient embeddings from the GAT layer

for all subjects in the cohort would be extracted. These embeddings would then be used as the input to train a VAE. This VAE would learn to project each patient into a continuous latent space, where their coordinates would represent their unique, personalized "progression signature," precisely as envisioned by the roadmap.¹

The scientific payoff of this addition would be immense. It would allow for a much more nuanced analysis of heterogeneity than discrete clustering, enabling researchers to investigate the biological drivers (e.g., specific genetic variants or proteomic profiles) that correlate with different axes of this continuous space. This directly addresses the roadmap's ultimate goal for this frontier: connecting the "what" of disease heterogeneity (the observable progression pattern) with the "why" (the underlying biological mechanism), thereby generating clear, testable hypotheses for future research.¹

Frontier 4 Alignment Analysis: The Prodromal Frontier

Roadmap Mandate: Predicting Onset to Enable Prevention

The fourth frontier represents the most significant strategic evolution of the PPMI study and its highest priority for future research.¹ The roadmap describes the massive investment in the recruitment and deep phenotyping of a large prodromal cohort as the "centerpiece of the PPMI expansion," reflecting a field-wide consensus that the ultimate goal must be to move from treatment to prevention.¹ To achieve this, researchers must be able to accurately identify individuals in the earliest, pre-symptomatic stages of the disease.

The primary goal of research on this frontier is therefore to build accurate, high-resolution risk models that can predict which at-risk individuals will go on to receive a clinical diagnosis of PD—a process known as phenoconversion.¹ The roadmap explicitly states that such a model is "the essential tool required for the design of future clinical trials for preventative therapies," as it would allow for the enrichment of trial cohorts with individuals who have a high probability of reaching the trial's endpoint, dramatically increasing statistical power and feasibility.¹ The mandate is unequivocal: the prodromal frontier is the most critical area for high-impact computational research using PPMI data.

GIMAN's Cohort Definition: Exclusive Focus on Manifest PD

The GIMAN research plan lays out a stringent, rule-based procedure for cohort selection to ensure the integrity of its analysis.³ The inclusion criteria for the PD cohort are explicitly limited to participants with an established clinical diagnosis of PD within two years of their baseline visit, who are

de novo (untreated with dopaminergic medications at baseline), and who have imaging evidence of a dopaminergic deficit confirmed by a baseline DAT-SPECT scan.³ The plan also defines a healthy control cohort. The document makes no mention of the PPMI prodromal cohort, and its inclusion criteria effectively exclude any individual who has not yet received a formal clinical diagnosis of PD.³

Assessment of Alignment: None. Critical Gap.

This represents the most significant and fundamental misalignment between the GIMAN research plan and the PPMI strategic roadmap. By focusing solely on modeling progression *after* diagnosis, the plan completely ignores the roadmap's highest-priority area for future research. It is, in effect, applying a technologically advanced methodology to a problem that the roadmap has clearly deprioritized in favor of the more pressing and impactful challenge of predicting disease onset.¹ This omission is not a minor gap but a fundamental divergence from the strategic direction of the entire PPMI enterprise.

A Method Perfectly Suited for a Problem It Ignores

The core conceptual strength of the GIMAN framework is its ability to model heterogeneity and inter-patient relationships within a cohort by constructing a patient similarity graph.³ The PPMI prodromal cohort is the epitome of a heterogeneous population. It is comprised of individuals with a diverse array of risk factors—including genetic variants (e.g.,

LRRK2, *GBA*), REM Sleep Behavior Disorder (RBD), and olfactory loss—who will inevitably follow vastly different trajectories.¹ Some will convert to clinical PD quickly, some will convert slowly, and a significant portion will remain stable, representing a cohort of "resilient non-converters".¹

The central challenge of predicting phenoconversion is therefore a problem of learning to separate these latent subgroups based on their baseline multimodal data. The GIMAN patient similarity graph is an ideal methodology for this exact problem.²⁶ A similarity graph could be constructed based on baseline risk factors (genetics, UPSIT scores, RBD screening questionnaires, DAT imaging metrics, etc.). The GNN could then learn to propagate information between similar at-risk individuals, naturally identifying communities of "likely converters" and distinguishing them from "resilient" individuals. This approach could uncover the multimodal signatures of neuroprotection or compensation that the roadmap highlights as a profound and valuable line of inquiry.¹ The GIMAN method is a key that is perfectly shaped to fit the lock of the prodromal frontier, yet the research plan does not even attempt to put it in the door.

Strategic Recommendation: Propose a Parallel GIMAN Model for the Prodromal Cohort

Addressing this critical gap requires a major revision and reframing of the entire research proposal. The project should be restructured to include two parallel, flagship analyses that showcase the versatility of the GIMAN framework across the full disease spectrum:

1. **GIMAN-Progression:** The currently proposed model, focused on predicting the dynamic progression of disability milestones in the *de novo* PD cohort (as revised per the recommendation for Frontier 1).
2. **GIMAN-Conversion:** A new, high-priority model focused on predicting the time-to-phenoconversion in the PPMI *prodromal* cohort.

This revision would transform the proposal from a niche methodological paper into a landmark study that directly addresses the number one strategic priority of the PPMI. It would demonstrate that the research team understands the foundation's mission and is applying its innovative methodology to the most pressing and impactful scientific question in the field today. A dual-model approach would not only double the potential impact of the research but would also serve as a powerful validation of the GIMAN framework's generalizability to different stages of the disease and different prediction tasks. This strategic pivot is essential for the project to be considered a top-tier, high-impact endeavor fully aligned with the PPMI's forward-looking vision.

Frontier 5 Alignment Analysis: Explainable AI and Causal Inference

Roadmap Mandate: Opening the Black Box and Probing Causality

The fifth frontier is grounded in the principle that for machine learning models to be truly useful in a clinical and scientific context, prediction is not sufficient; we must understand *why* a model makes a particular prediction.¹ The roadmap calls for the integration of Explainable AI (XAI) techniques to transform predictive models from opaque "black boxes" into powerful, transparent engines for novel biomarker discovery.¹ It specifically suggests methods like SHAP (SHapley Additive exPlanations) for feature importance ranking and saliency maps for interpreting neuroimaging models.¹ Looking further ahead, the roadmap also encourages the exploration of Causal Inference frameworks, such as Mendelian Randomization or Causal Discovery algorithms, to begin moving from identifying correlations to generating plausible causal hypotheses about disease mechanisms.¹

GIMAN's Interpretability Framework: A Multi-Scale Deep Dive

The GIMAN research plan features a comprehensive and exceptionally well-structured multi-scale interpretability framework as its fourth specific aim.³ This demonstrates a deep commitment to transparency and scientific discovery that goes far beyond a superficial nod to XAI. The plan correctly identifies that a hierarchical model like GIMAN requires a multi-faceted interpretation strategy and proposes specific, appropriate methods to probe the model at three distinct levels of analysis³:

- **Population Level:** GNNExplainer will be applied to the trained GAT layer to identify influential patient subgraphs and the most critical inter-patient similarities, thereby discovering and characterizing novel, data-driven patient archetypes.³
- **Individual Patient Level:** SHAP will be applied to the full model's output to compute the contribution of each input feature to an individual patient's prediction. This enables the generation of personalized biomarker rankings and provides the foundation for a trustworthy clinical decision support tool.³
- **Biomarker Level:** For the imaging encoder, Gradient-weighted Class Activation Mapping (Grad-CAM) will be used to produce saliency maps that highlight the specific brain regions the model focused on, allowing for validation of biological plausibility and the discovery of novel neuroanatomical correlates of progression.³

Assessment of Alignment: Strong

The GIMAN plan's commitment to XAI is a major strength and is perfectly aligned with the biomarker discovery portion of this frontier. The plan does not merely list XAI techniques; it articulates a clear strategy for how these methods will be deployed at different scales to translate the model's internal workings into meaningful biological and clinical insights. This approach directly fulfills the roadmap's vision of using XAI as an engine for hypothesis generation. The plan does not, however, explicitly address the roadmap's more forward-looking suggestion to explore causal inference.

The Graph as a Latent Causal Structure

The GIMAN framework, through its execution, naturally generates a rich set of outputs that are prime targets for future causal analysis. The patient similarity graph itself, where edges represent complex, multi-domain relationships, can be interpreted as a data-driven map of phenotypic relationships. The results from GNNExplainer, which identify the most important edges and subgraphs for prediction, could highlight potential pathways of influence within this map. This learned graph structure provides a powerful, pre-processed, and regularized input for the very causal discovery algorithms (e.g., PC algorithm, FCI) that the roadmap suggests exploring.¹ The GIMAN framework, therefore, does more than just predict outcomes; it generates a rich relational structure that can serve as a foundational artifact for the next wave of inquiry into the causal mechanisms of PD progression.

Strategic Recommendation: Frame Causal Inference as a Key Future Direction

While adding a full aim dedicated to causal inference may be beyond the scope of the initial project, the "Limitations and Future Directions" section of the GIMAN plan should be significantly expanded to explicitly incorporate this concept.³ The plan should propose that a key next step following the successful completion of the proposed aims will be

to apply causal discovery algorithms to the learned GIMAN graph structure and the

top-ranked predictive biomarkers identified by SHAP analysis. The goal of this future work would be to generate a plausible causal graph of PD progression, moving from the "what" of prediction to the "why" of mechanism.

This addition would demonstrate strategic foresight, showing that the research team is not only addressing the current state-of-the-art in XAI but is also positioning their work as a critical foundation for the next frontier of computational neurology. This would achieve full alignment with the entire scope of Frontier 5, from immediate biomarker discovery to long-term causal understanding.

Consolidated Gap Analysis and Strategic Recommendations

Synthesis of Findings

The comprehensive, frontier-by-frontier analysis reveals that the GIMAN research plan proposes a technically powerful and innovative framework that is, in its current form, strategically under-leveraged. The core architectural innovations—particularly the use of GNNs for population modeling and the gradual fusion pipeline—are highly commendable and align perfectly with the methodological advancements called for in the PPMI roadmap. However, the plan's strategic vision is constrained by a narrow focus on a single cohort (de novo PD) and a limited set of static prognostic endpoints. This narrow application causes it to miss significant opportunities in dynamic progression modeling, biological grounding with the SAA biomarker, continuous heterogeneity analysis, and—most critically—the high-priority prodromal frontier. The following matrix consolidates the identified gaps and provides actionable, strategic recommendations to bridge them.

Strategic Alignment Matrix

Strategic Frontier (from PPMI)	GIMAN Plan	Identified Gap &	Strategic
-----------------------------------	------------	------------------	-----------

Roadmap) ¹	Alignment ³	Missed Opportunity	Recommendation
1. Dynamic Progression Modeling	Partial. Proposes predicting the slope of MDS-UPDRS Part III and binary cognitive conversion over a fixed 36-month window. ³	The endpoints are static summaries, not dynamic trajectories. This misses the opportunity for time-to-event forecasting of clinically meaningful disability milestones. ¹	Revise Aim 2 to incorporate deep survival analysis (e.g., DeepSurv, DeepHit) to predict the individualized time-to-event for key disability milestones identified in the PPMI cohort. ⁵
2. Deep Multi-Modal Fusion	Strong. Employs a state-of-the-art "gradual fusion" architecture with modality-specific encoders, a Graph Attention Network, and cross-modal attention. ³	The plan does not incorporate the alpha-synuclein Seed Amplification Assay (SAA) biomarker , which the roadmap identifies as a revolutionary "biological ground truth". ¹	Add a secondary objective to Aim 2: train the GIMAN model to predict a patient's binary SAA status from non-invasive and minimally invasive data, creating a valuable proxy biomarker for synucleinopathy.
3. Deconstructing Heterogeneity	Partial. Addresses heterogeneity implicitly through post-hoc analysis of the learned graph structure to identify "patient archetypes". ³	Lacks a primary, unsupervised learning objective. Misses the roadmap's call for modeling heterogeneity as a continuous "progression signature" in a latent space. ¹	Add a new Aim to apply a Variational Autoencoder (VAE) to the final GIMAN-generated patient embeddings to explicitly model the continuous latent space of disease heterogeneity. ²¹
4. The Prodromal	None. The plan's	This represents a	Major Revision:

Frontier	cohort definition is strictly limited to <i>de novo</i> PD patients and healthy controls, explicitly excluding the prodromal population. ³	complete and critical gap, ignoring the roadmap's #1 strategic priority: predicting phenoconversion in the prodromal cohort to enable future prevention trials. ¹	Propose a parallel GIMAN-Conversion model dedicated to predicting the time-to-phenoconversion in the PPMI prodromal cohort, elevating the project's strategic importance.
5. XAI & Causal Inference	Strong. Includes a comprehensive, multi-scale XAI framework for biomarker discovery, using GNNExplainer, SHAP, and Grad-CAM. ³	The plan fully addresses the XAI goal but does not engage with the forward-looking suggestion to explore causal inference as a future direction. ¹	Expand the "Future Directions" section to explicitly propose applying causal discovery algorithms to the learned graph structure and identified biomarkers as a key next step in the research program.

Conclusion: A Revised Path to High-Impact Research

Final Assessment

The GIMAN research plan, in its current form, presents a significant technical advance in the field of computational neurology. Its novel graph-based architecture for population-informed prediction and its deep commitment to a multi-scale interpretability framework are laudable strengths that set it apart from conventional prognostic models. The methodological rigor and technical sophistication detailed in the plan are of a high caliber.

The Call to Action

However, this analysis concludes that the plan's strategic application does not match its technical ambition. By focusing on a narrow slice of the PD spectrum and employing static endpoints, the project misses multiple opportunities to address the most pressing and impactful questions outlined in the PPMI strategic roadmap. The recommendations provided in this report—embracing time-to-event analysis, incorporating the SAA biomarker, explicitly modeling continuous heterogeneity, and, most critically, tackling the prodromal frontier—are not merely suggestions for incremental improvement. They represent a call to action to strategically pivot the project, aligning its powerful methodology with the highest-priority goals of the PD research community.

Vision for the Future

By incorporating these strategic revisions, the GIMAN project can be elevated from a high-quality methodological study into a landmark investigation with the potential for transformative impact. The revised project would represent a comprehensive, dual-pronged program that uses a unified, powerful methodology to address the two most critical questions in PD research today: "How does the disease progress after diagnosis?" and "How can we predict who will get the disease in the first place?" This revised plan would not only constitute a more impactful investment of research resources but would also stand as a model for how to fully leverage the immense, multidimensional potential of the PPMI data ecosystem. In doing so, it would perfectly realize the ambitious vision of the strategic roadmap, paving the way for a future of more precise, personalized, and preventative neurological care.

Works cited

1. PPMI Data Research Next Steps.pdf
2. (PDF) Machine learning within the Parkinson's progression markers initiative: Review of the current state of affairs - ResearchGate, accessed October 5, 2025, https://www.researchgate.net/publication/368468281_Machine_learning_within_the_Parkinson's_progression_markers_initiative_Review_of_the_current_state_of_affairs
3. Addressing Reviewer Concerns for Manuscript Revision.pdf
4. Graph Neural Networks for Parkinson's Disease Monitoring and Alerting - PubMed Central, accessed October 5, 2025, <https://pmc.ncbi.nlm.nih.gov/articles/PMC10648881/>

5. Parkinson's Progression Markers Initiative: A Milestone-Based ..., accessed October 5, 2025, <https://pubmed.ncbi.nlm.nih.gov/37458046/>
6. Deep Learning Models for Predicting Clinical Outcomes in Parkinson's Diseases - JETIR.org, accessed October 5, 2025, <https://www.jetir.org/papers/JETIR2410325.pdf>
7. Prediction of conversion to Alzheimer's disease using deep survival ..., accessed October 5, 2025, <https://pmc.ncbi.nlm.nih.gov/articles/PMC7425528/>
8. (PDF) Time-to-event prediction using survival analysis methods for Alzheimer's disease progression - ResearchGate, accessed October 5, 2025, https://www.researchgate.net/publication/357472854_Time-to-event_prediction_using_survival_analysis_methods_for_Alzheimer's_disease_progression
9. Architecture of DeepSurv model used in this study The architecture is the same as in [9]. - ResearchGate, accessed October 5, 2025, https://www.researchgate.net/figure/Architecture-of-DeepSurv-model-used-in-this-study-The-architecture-is-the-same-as-in-9_fig1_377597013
10. Deep Neural Networks for Predicting Recurrence and Survival in Patients with Esophageal Cancer After Surgery - arXiv, accessed October 5, 2025, <https://arxiv.org/html/2409.00163v1>
11. DeepHit: A Deep Learning Approach to Survival Analysis With Competing Risks | Request PDF - ResearchGate, accessed October 5, 2025, https://www.researchgate.net/publication/361527038_DeepHit_A_Deep_Learning_Approach_to_Survival_Analysis_With_Competing_Risks
12. (PDF) Parkinson's Progression Markers Initiative: A Milestone-Based Strategy to Monitor PD Progression - ResearchGate, accessed October 5, 2025, https://www.researchgate.net/publication/371027579_Parkinson's_Progression_Markers_Initiative_A_Milestone-Based_Strategy_to_Monitor_PD_Progression
13. Parkinson's Progression Markers Initiative: A Milestone-Based ..., accessed October 5, 2025, <https://pmc.ncbi.nlm.nih.gov/articles/PMC10578214/>
14. Alpha-Synuclein Seed Amplification Assays in Parkinson's Disease: A Systematic Review and Network Meta-Analysis - MDPI, accessed October 5, 2025, <https://www.mdpi.com/2039-7283/15/6/107>
15. Assessment of heterogeneity and disease onset in the Parkinson's Progression Markers Initiative (PPMI) cohort using the α -synuclein seed amplification assay: a cross-sectional study, accessed October 5, 2025, <https://parkinsonsroadmap.org/report/assessment-of-heterogeneity-and-disease-onset-in-the-parkinsons-progression-markers-initiative-ppmi-cohort-using-the-%CE%B1-synuclein-seed-amplification-assay-a-cross-sectional-study/>
16. Machine learning models for diagnosis and prognosis of Parkinson's disease using brain imaging: general overview, main challenges, and future directions - ResearchGate, accessed October 5, 2025, https://www.researchgate.net/publication/372508543_Machine_learning_models_for_diagnosis_and_prognosis_of_Parkinson's_disease_using_brain_imaging_general_overview_main_challenges_and_future_directions
17. Baseline α -synuclein seeding activity and disease progression in sporadic and genetic Parkinson's disease in the PPMI cohort - PMC - PubMed Central,

- accessed October 5, 2025, <https://pmc.ncbi.nlm.nih.gov/articles/PMC12354789/>
18. α -Synuclein Seed Amplification Assay Amplification Parameters and the Risk of Progression in Prodromal Parkinson Disease | Neurology, accessed October 5, 2025, <https://www.neurology.org/doi/10.1212/WNL.0000000000210279>
 19. Predicting CSF α -Synuclein Seed Amplification Assay Status From Demographics and Clinical Data | Neurology Open Access, accessed October 5, 2025, <https://www.neurology.org/doi/10.1212/WN9.0000000000000008>
 20. Access Data | Parkinson's Progression Markers Initiative, accessed October 5, 2025, <https://www.ppmi-info.org/access-data-specimens/download-data>
 21. Novel multi-omics deconfounding variational autoencoders can ..., accessed October 5, 2025, <https://pmc.ncbi.nlm.nih.gov/articles/PMC11483139/>
 22. [2303.12706] Multi-modal Variational Autoencoders for normative modelling across multiple imaging modalities - arXiv, accessed October 5, 2025, <https://arxiv.org/abs/2303.12706>
 23. Variational autoencoder provides proof of concept that compressing CDT to extremely low-dimensional space retains its ability of distinguishing dementia. - National Genomics Data Center (CNCB-NGDC), accessed October 5, 2025, <https://ngdc.cncb.ac.cn/openlb/publication/OLB-PM-35568709>
 24. Predicting Phenoconversion in Isolated RBD: Machine Learning and ..., accessed October 5, 2025, <https://pmc.ncbi.nlm.nih.gov/articles/PMC12015906/>
 25. Study Cohorts | Parkinson's Progression Markers Initiative, accessed October 5, 2025, <https://www.ppmi-info.org/study-design/study-cohorts>
 26. Graph Neural Networks for Disease Prediction Based on Microbial Profiles - Journal of Artificial Intelligence Research and Applications, accessed October 5, 2025, <https://www.e-jara.org/journal/view.html?uid=1527&vmd=Full>
 27. (PDF) Graph Neural Networks for Modeling Disease Relationships: A Framework for Multi-Disease Diagnostics and Comorbidity Prediction - ResearchGate, accessed October 5, 2025, https://www.researchgate.net/publication/386460895_Graph_Neural_Networks_for_Modeling_Disease_Relationships_A_Framework_for_Multi-Disease_Diagnostics_and_Comorbidity_Prediction
 28. Graph Neural Networks for Heart Failure Prediction on an EHR-Based Patient Similarity Graph - arXiv, accessed October 5, 2025, <https://arxiv.org/html/2411.19742v1>