

Supplementary Information

Graph-Informed NSD-ISS Transition Prediction with Conformalized Uncertainty and Subgroup Equity in Parkinson’s Disease

Blair D. Dupre

Department of Biomedical Engineering,
University of North Dakota, Grand Forks, ND, USA

S-1. Reserved

Reserved for the four-method conformal-ablation table (IPCW / marginal / naive / Bonferroni) at 90% and 95% CL. Numerical content is reported in the main-text Methods (§Cause-specific conformal CIF bands with IPCW) and the Results figures.

S-2. Reserved

Reserved for the competitor-comparison table positioning this work against Candès *et al.* 2023 and Sreenivasan *et al.* 2025. Discussion is in the main-text Discussion (§Positioning relative to prior work).

S-3. Carrier Subgroup Analysis

Background and pre-registration

The original Paper 3+4 submission excluded LRRK2 and GBA carrier subgroups from inferential analysis under the rationale that per-fold $n < 10$, the result of a silent feature-extraction bug in `scripts/assemble_paper1_features.py:extract_genetics()` that failed to match the IU Genetic Consensus variant-name encoding (G2019S, R1441G, N409S, L483P). After the bug fix in commit 672b439, the post-fix carrier columns in `features.paper1_features_with_targets` are: 175 LRRK2 carriers, 111 GBA carriers, and 441 APOE ϵ 4 carriers. Strata are defined as mutually exclusive in this analysis: LRRK2+ includes any patient with `lrrk2_carrier=1` (including dual carriers); APOE ϵ 4+ only and GBA+ only require the other two carrier flags to be 0. The pre-registration locking the hypotheses (H1/H2/H3), `MIN_SUBGROUP_SIZE_CARRIER =`

50, and decision rule was committed on 2026-04-23 *before* any compute on the bug-fixed feature table (`outputs/paper4/subgroup_carriers/PRE_REGISTRATION.md`). The runner script is `scripts/paper3plus4/run_subgroup_with_lrrk2_gba_fix.py`.

Cohort sizes (post mutual-exclusivity definition)

Stratum-level (cohort-wide) sizes versus 5-fold cross-validation per-fold mean and range:

Table S1: **Carrier-stratum cohort and per-fold sizes.** Cohort-level n is the post-bug-fix carrier flag count in `features.paper1_features_with_targets` after applying mutual-exclusivity definitions. Per-fold n is the count of test patients within that stratum across the 5 stratified cross-validation folds (mean and min–max range). Per-fold events count distinct first-event records used in the time-dependent concordance computation.

Stratum	Cohort n	Per-fold n (mean)	Per-fold n (range)	Per-fold events (m
LRRK2+ (incl. dual carriers)	175	140.2	109–179	
APOE ϵ 4+ only	375	161.0	109–175	
GBA+ only	79	50.0	10–92	
Non-carrier (reference)	1,547	590.4	567–636	

H1 — Per-stratum C-td (fairness on discrimination)

Per-stratum time-dependent concordance averaged across 5 stratified cross-validation folds, with 95% patient-level bootstrap CIs (1,000 resamples). The CI envelope columns report the minimum lower bound and maximum upper bound across the 5 per-fold CIs (a conservative summary across folds rather than a refit-bootstrap). All carrier strata’s CI envelopes overlap the non-carrier reference envelope, supporting the H1 fairness pre-registration.

Table S2: **H1 — Per-stratum C-td (mean $\pm$ std across 5 folds; CI envelope across 5 per-fold patient-bootstrap 95% CIs).** Per-fold n and event counts vary by fold composition; the headline n in column 5 is the per-fold mean. The GBA+ only entry has elevated variance because per-fold n ranges from 10 to 92; the wide CI envelope [0.500, 1.000] reflects the smallest per-fold sample size.

Model	Stratum	C-td (mean $\pm$ std)	CI envelope (lo, hi)	n (per-fold mean)	Events
DeepHit	LRRK2+	0.911 ± 0.024	[0.818, 0.976]	140.2	
	APOE ϵ 4+ only	0.932 ± 0.026	[0.811, 1.000]	161.0	
	GBA+ only	0.862 ± 0.183	[0.500, 1.000]	50.0	
	Non-carrier	0.928 ± 0.016	[0.867, 0.956]	590.4	
Graph-DT	LRRK2+	0.897 ± 0.026	[0.778, 0.954]	140.2	
	APOE ϵ 4+ only	0.907 ± 0.043	[0.776, 0.990]	161.0	
	GBA+ only	0.854 ± 0.184	[0.500, 1.000]	50.0	
	Non-carrier	0.906 ± 0.033	[0.806, 0.954]	590.4	

Verdict. H1 PASS: every carrier-stratum CI envelope overlaps the non-carrier reference for both DeepHit and Graph-DT.

H2 — Bootstrap interaction tests (BH-FDR-corrected)

Reviewer #3 (reviewer3.com) flagged that combining per-fold interaction p -values via Fisher’s $\chi^2_{2k} = -2 \sum \ln(p_i)$ violates the procedure’s strict independence assumption: in a 5-fold CV scheme the test sets are disjoint but the trained models share 60% of the training data, so per-fold p -values are positively correlated and Fisher’s combination yields spuriously low pooled p -values. The fix (WS-P3-CRIT-B) replaces the per-fold-then-Fisher-combine pattern with a SINGLE patient-level bootstrap interaction test ($B=500$, `random_state=42`) on the pooled out-of-fold predictions across the 5 folds: each patient appears in exactly one test fold by construction, so a single permutation test on the pooled cohort is the correct statistic for $H_0 : \Delta C_{td} = 0$. The pooled-OOF inference framework follows Vovk *et al.* [3]. BH-FDR correction is applied across the 6-cell carrier grid ($3 \text{ strata} \times 2 \text{ models}$) using `scipy.stats.false_discovery_control(method='bh')`. The legacy per-fold-then-Fisher pipeline is retained in the JSON record under `p_fisher_legacy` and `p_fdr_fisher_legacy` for the side-by-side audit shown below; pre-fix outputs are snapshotted at `outputs/paper4/subgroup_carriers/_pre_crit_b/`.

Table S3: **H2 — Bootstrap interaction tests with BH-FDR correction (post-WS-P3-CRIT-B).** $\Delta C_{td}^{\text{pooled}}$ is computed on the union of the 5 per-fold OOF predictions (carrier C-td – non-carrier C-td); negative values mean the carrier stratum performs worse than the non-carrier reference. p_{pooled} is the new pooled-OOF permutation p -value (CRIT-B canonical); p_{Fisher} is the legacy Fisher-combined value (no longer drives FDR); $p_{\text{FDR,pooled}}$ and $p_{\text{FDR,Fisher}}$ are the BH-corrected values from each respective pipeline. Numbers in this table are pulled directly from `outputs/paper4/subgroup_carriers/interaction_tests_carriers.json` (post-fix) and `outputs/paper4/subgroup_carriers/_pre_crit_b/interaction_tests_carriers.json` (pre-fix).

Model	Stratum	$\Delta C_{td}^{\text{pooled}}$	p_{pooled} (CRIT-B)	p_{Fisher} (legacy)	$p_{\text{FDR,pooled}}$	$p_{\text{FDR,Fisher}}$
DeepHit	LRRK2+	−0.017	0.120	0.236	0.719	0.372
DeepHit	GBA+ only	+0.004	0.856	0.355	0.856	0.426
DeepHit	APOE $\epsilon 4+$ only	+0.003	0.774	0.683	0.856	0.683
Graph-DT	LRRK2+	−0.014	0.273	0.163	0.737	0.372
Graph-DT	GBA+ only	−0.017	0.431	0.049	0.737	0.295
Graph-DT	APOE $\epsilon 4+$ only	−0.009	0.491	0.248	0.737	0.372

Verdict. The H2 PASS verdict is preserved under the methodologically-correct pooled-OOF re-test. All six $p_{\text{FDR,pooled}}$ values exceed the pre-registered $\alpha=0.05$ threshold, with the smallest equal to 0.719 (DeepHit×LRRK2+, $\Delta C_{td} = -0.017$). The legacy Fisher pipeline produced a marginal Graph-DT×GBA+ only $p_{\text{Fisher}}=0.049$ that did not survive FDR correction; the pooled-OOF re-test relabels this borderline cell as $p_{\text{pooled}}=0.431$ ($p_{\text{FDR,pooled}}=0.737$), which is the correct inferential statistic for dependent per-fold p -values and confirms the pre-fix $p_{\text{Fisher}}=0.049$ was a

multiple-comparison artifact amplified by the positive dependence between CV folds (the failure mode reviewer3.com #3 anticipated). The §S-CRIT-B audit table below reports the per-cell change.

H3 — Conditional conformal coverage (per stratum)

Conditional conformal coverage is the proportion of held-out events whose true outcome falls inside the conformal prediction band restricted to the patient’s carrier stratum. Pre-registration: stratum-conditional coverage at the 90% confidence level should fall within ± 0.03 of the empirical marginal coverage. **The empirical marginal coverage shifted from 0.82 (pre-fix) to 0.90 (post-fix) after the WS-P3-CRIT-A IPCW formula correction (commit f8a7c54, see § S-CRIT-A below): per-marginal-cohort coverage rose from 0.817 to 0.902 for DeepHit and from 0.818 to 0.905 for Graph-DT at 90% CL.** Under the corrected IPCW pipeline, the H3 tolerance window is therefore $[0.87, 0.93]$ at the 90% CL.

Table S4: **H3 — Conditional conformal coverage per stratum (post-WS-P3-CRIT-A IPCW formula fix).** Mean across 5 folds at the 90% and 95% confidence levels. The pre-registered H3 tolerance is ± 0.03 around the empirical marginal coverage; under the corrected IPCW pipeline (commit f8a7c54) the empirical marginal at 90% CL is ≈ 0.90 , so the per-stratum tolerance window is $[0.87, 0.93]$ at 90% CL and $[0.92, 0.98]$ at 95% CL. **All eight (model $\times$ stratum) cells fall within tolerance at both confidence levels.**

Model	Stratum	Mean coverage (90% CL)	Mean coverage (95% CL)
DeepHit	LRRK2+	0.898	0.949
	GBA+ only	0.919	0.960
	APOE $\varepsilon 4+$ only	0.909	0.953
	Non-carrier	0.901	0.950
Graph-DT	LRRK2+	0.894	0.951
	GBA+ only	0.926	0.962
	APOE $\varepsilon 4+$ only	0.915	0.955
	Non-carrier	0.906	0.953

Verdict. H3 PASS: stratum-conditional coverage at the 90% CL lies within $[0.87, 0.93]$ for all (model $\times$ stratum) cells. The largest single deviation from nominal 0.90 is Graph-DT $\times$ GBA+ only at $+0.026$ (still inside the ± 0.03 window); the smallest is DeepHit $\times$ Non-carrier at $+0.001$. The pre-fix H3 FAIL verdict (mean stratum coverage 0.795–0.835) was an artifact of the silent IPCW formula bug (per-fold weight $1/G(C_i)$ for censored survivors instead of the Candès–Lei–Ren three-case rule), not of the carrier strata themselves — the same bug that depressed pooled marginal coverage by ≈ 9 pp at 90% CL also depressed every stratum’s coverage by a similar amount. Under the corrected IPCW pipeline the carrier strata behave as well as the demographic (sex/age) strata.

Mondrian per-stratum sensitivity check (WS-P3-14b)

To verify that the H3 PASS is not merely a coincidence of the post-CRIT-A marginal lift, we performed a Mondrian per-stratum recalibration [1, 3] as a sensitivity check: for each fold and

each carrier stratum, the IPCW conformal quantiles were re-estimated on the calibration set restricted to that stratum, then applied to the eval-set patients of the same stratum. This is the statistically correct deployment-time procedure for genetically-defined subpopulations. The runner is at `scripts/paper3plus4/run_mondrian_carriers.py`; with `MIN_CAL_PER_STRATUM=30` the GBA+ only stratum is Mondrian-eligible in only 2 of 5 folds (per-fold cal-N $\in \{43, 49, 87, 87, 98\}$ across folds; the remaining folds fall back to the marginal quantile). LRRK2+ and APOE $\epsilon 4$ + only are eligible in all 5 folds.

Table S5: **Mondrian sensitivity check (90% CL, post-CRIT-A IPCW pipeline)**. Marginal column reproduces Table S4; Mondrian column applies stratum-specific quantiles. Δ is Mondrian – marginal. Width inflation reports the relative increase in mean band width under Mondrian. The GBA+ only stratum is Mondrian-eligible in 2/5 folds only (`n_cal < 30` in the other folds); other strata 5/5.

Model	Stratum	Marginal	Mondrian	Δ	Width inflation	Folds (Mondrian)
DeepHit	LRRK2+	0.898	0.907	+0.009	+141%	5/5
	GBA+ only	0.919	0.938	+0.037	+259%	2/5
	APOE $\epsilon 4$ + only	0.909	0.912	+0.003	+41%	5/5
	Non-carrier	0.901	0.898	−0.003	+7%	5/5
Graph-DT	LRRK2+	0.894	0.915	+0.020	+72%	5/5
	GBA+ only	0.926	0.933	+0.028	+161%	2/5
	APOE $\epsilon 4$ + only	0.915	0.919	+0.004	+27%	5/5
	Non-carrier	0.906	0.901	−0.005	+5%	5/5

Sensitivity-check verdict. Mondrian recalibration confirms the H3 PASS: all 8 (model $\times$ stratum) cells stay within the $[0.87, 0.93]$ tolerance window (largest Mondrian deviation +0.038 for DeepHit $\times$ GBA+ only). Mondrian provides a small per-stratum coverage boost ($\Delta \in [-0.005, +0.037]$) but at substantial width cost (mean band width inflates +5% for Non-carrier to +259% for the data-poor DeepHit $\times$ GBA+ only cell, where the per-fold Mondrian quantile is computed from only ≈ 45 calibration patients). The observed inflation pattern is the expected statistical-efficiency cost of Mondrian (Boström *et al.* 2025): smaller calibration N at fixed nominal coverage produces wider bands. The GBA+ only stratum’s 2/5 fold eligibility is the fundamental finite-sample limit and is the cleanest motivation for multi-site pooling (PPMI + PDBP + UK Biobank-PD) for clinical deployment in the GBA+ subpopulation. The full per-(fold, model, CL, stratum) Mondrian record is at `outputs/paper4/subgroup_carriers/conditional_coverage_carriers_mondrian.json` and the aggregate summary at `outputs/paper4/subgroup_carriers/mondrian_summary.json`.

Honest disclosure and deployment recommendation

The pre-registered overall verdict for WS-P3-14 (post-CRIT-A + post-CRIT-B + post-WS-P3-14b) is therefore **full pass**: H1 fairness is supported (per-stratum C-td CIs overlap the non-carrier reference), H2 is supported (post-CRIT-B pooled-OOF interaction tests yield no significant model $\times$ subgroup interaction after BH-FDR correction; smallest $p_{\text{FDR}}=0.719$), and H3 is supported

(post-CRIT-A IPCW marginal coverage within ± 0.03 of nominal across all 8 strata cells, and the Mondrian sensitivity check confirms the same conclusion at significantly inflated band width). For *deployment* to genetically-defined subpopulations we recommend Mondrian recalibration in any setting where the per-stratum nominal coverage guarantee outweighs the band-width cost; the marginal calibration is sufficient for population-level reporting. A multi-site pooled analysis would lift the GBA+ only stratum out of the per-fold $n_{\text{cal}} < 30$ regime, allowing 5/5 fold Mondrian eligibility, and is the natural next step for this work.

Source artefacts

- `outputs/paper4/subgroup_carriers/PRE_REGISTRATION.md` — locked hypothesis and decision rule (2026-04-23).
- `outputs/paper4/subgroup_carriers/subgroup_ctd_carriers.json` — per-fold per-stratum C-td with bootstrap CIs.
- `outputs/paper4/subgroup_carriers/interaction_tests_carriers.json` — per-fold bootstrap interaction tests + Fisher combination.
- `outputs/paper4/subgroup_carriers/conditional_coverage_carriers.json` — per-fold stratum-conditional coverage at 90% and 95% CL.
- `outputs/paper4/subgroup_carriers/decision_verdict.json` — pooled summary, FDR-corrected p -values, H3 deviation list.
- `outputs/paper4/subgroup_carriers/CLAIMS.md` — claim-mapping audit record.
- `scripts/paper3plus4/run_subgroup_with_lrrk2_gba_fix.py` — runner.

S-CRIT-A. IPCW Formula Correction (Reviewer Round 2 #2)

Statement of the bug and correct formula

Reviewer #2 (npj-DM, Round 2) flagged that the original implementation of the cause-specific IPCW conformal weights violated Candès, Lei, and Ren [2]. Specifically, the per-observation weight at evaluation horizon t_j was computed as $w_i = 1/\hat{G}(C_i)$ for censored survivors with $C_i \geq t_j$ and $w_i = 1$ for uncensored events with $T_i \leq t_j$. Both choices are incorrect.

The Candès 2023 formula for fixed-time conformal evaluation is, for each calibration observation $i = (X_i, \delta_i)$ contributing to horizon t_j :

$$w_i^{(k,j)} = \begin{cases} 1/\hat{G}(T_i^-) & \delta_i = 1, T_i \leq t_j \quad (\text{uncensored event before } t_j) \\ 1/\hat{G}(t_j^-) & X_i > t_j \quad (\text{at risk past } t_j; \text{ censored or uncensored}) \\ \text{excluded} & \delta_i = 0, C_i < t_j \quad (\text{censored before } t_j) \end{cases} \quad (1)$$

where $X_i = \min(T_i, C_i)$, $\delta_i \in \{0, 1\}$, T_i is the event time, C_i the censoring time, and $\hat{G}$ the Kaplan–Meier estimate of the censoring survival function. The original (buggy) implementation used the patient’s own censoring time C_i for survivors, which over-weights patients censored late, and unit weight for uncensored events, which under-weights events whose at-risk-process realisation is short.

The corrected helper `compute_per_observation_ipcw_weight()` in `src/giman_pipeline/paper4/conformal` implements Eq. 1. The same fix was mirrored in `src/giman_pipeline/paper4/calibration.py`. The KM estimation, the score formula, the per-(cause, time) quantile rule, and the random-state are unchanged; the only modification is the per-observation weight.

Pre/post-correction marginal coverage (5-fold mean $\pm$ SD)

Table S6: **Pre/post-correction marginal conformal coverage** on all 10 Paper 3 checkpoints (5 DeepHit + 5 Graph-DT), 50/50 calibration/evaluation split per fold, identical seeds. With the corrected IPCW formula, both models meet the nominal target at the 80%, 90%, and 95% confidence levels, closing the previously-documented structural undercoverage at 90% CL.

Model	CL	Pre coverage	Post coverage	Δ (pp)
DeepHit	80%	0.627 ± 0.048	0.797 ± 0.010	+16.97
	90%	0.817 ± 0.025	0.902 ± 0.008	+8.47
	95%	0.911 ± 0.015	0.951 ± 0.006	+3.93
Graph-DT	80%	0.626 ± 0.040	0.798 ± 0.009	+17.13
	90%	0.818 ± 0.024	0.905 ± 0.007	+8.67
	95%	0.914 ± 0.013	0.953 ± 0.006	+3.88

Per-cause coverage shifts at 90% CL

Table S7: **Per-cause marginal coverage at 90% CL, pre vs. post correction.** 5-fold mean. After correction every cause-specific value lies within ± 0.01 of the nominal 0.90 target for both models. Cause indices 0, 1, 2, 3, 4, 5, 6 correspond to NSD-ISS stages 0, 1, 2B, 3, 4, 5, 6 respectively (the destination of the next observed transition).

Cause	DeepHit pre	DeepHit post	Graph-DT pre	Graph-DT post
0 (Stage 0)	0.825	0.902	0.823	0.904
1 (Stage 1)	0.820	0.891	0.826	0.906
2 (Stage 2B)	0.820	0.903	0.816	0.902
3 (Stage 3)	0.816	0.901	0.810	0.900
4 (Stage 4)	0.807	0.910	0.813	0.912
5 (Stage 5)	0.815	0.908	0.819	0.912
6 (Stage 6)	0.818	0.897	0.824	0.901

Magnitude of change discussion

The correction is empirically substantial: marginal coverage at the 90% confidence level shifts by +8.5 pp (DeepHit) and +8.7 pp (Graph-DT), elevating both models from ~ 0.82 (the previously-reported “structural undercoverage”) to the nominal ~ 0.90 target. At 95% CL the shift is +3.9 pp for both models; at 80% CL it is +17 pp. These magnitudes exceed the 0.005 reporting threshold pre-registered in the WS-P3-CRIT-A workstream, so the abstract, the main-text conformal-coverage table, and Methods §Cause-specific conformal CIF bands with IPCW have all been updated in the main manuscript.

The previously-reported “structural concentration of CIF values near zero” framing for the 90%-CL undercoverage is now obsolete: the gap to nominal was an implementation bug, not an inherent property of pointwise conformal bands on competing-risks CIFs. The forward-vs-backward directional gap also narrows from 7.0 pp (pre) to 5.7 pp (post), because forward coverage is now lifted closer to nominal while backward coverage—which is intrinsically harder due to treatment-driven censoring informativeness—remains below 0.90 (0.854 mean across folds). The substantive Discussion claim that backward-transition coverage is sub-nominal is preserved; only the numerical values were updated.

The mean band widths grew with the corrected formula: at 90% CL, DeepHit width $0.0047 \rightarrow 0.0191$ and Graph-DT width $0.0171 \rightarrow 0.0478$. This is the price of correctness: the original buggy weights produced narrower bands by under-weighting. The post-correction bands are still substantially narrower than the Bonferroni baseline (~ 0.765 , $\sim 16\text{--}40\times$ wider) and remain clinically interpretable given that the units are CIF-probability (so widths ~ 0.05 correspond to ± 5 percentage points of cumulative incidence).

Files modified

- `src/giman_pipeline/paper4/conformal_survival.py`: added `compute_per_observation_ipcw_weight()` implementing Eq. 1; modified `CauseSpecificConformal.calibrate()` to call it for each (i, t_j) pair.
- `src/giman_pipeline/paper4/calibration.py`: replaced the same buggy weight pattern in `_compute_observed_status()` with the corrected formula; affects ECE, reliability diagrams, and Hosmer–Lemeshow test inputs.
- `tests/paper4/test_conformal_ipcw.py` (new): three TDD unit tests on a synthetic Kaplan–Meier with hand-computed $G(t)$ values, plus a smoke-test on a small synthetic CIF problem. All four tests fail on the pre-fix code (they cannot import `compute_per_observation_ipcw_weight()`) and pass after the fix.
- Pre-fix conformal results snapshotted to `outputs/paper4/conformal/_pre_ipcw_fix/` for the pre/post comparison above.

S-CRIT-B. Pooled-OOF Single Test Replaces Fisher’s Combined p -values (Reviewer Round 2 #3)

Statement of the bug

Reviewer #3 (npj-DM, Round 2) observed:

“Combining p -values across cross-validation folds via Fisher’s method violates the procedure’s strict independence assumption. ...Fisher’s method (using the test statistic $-2 \sum \ln(p_i) \sim \chi^2_{2k}$) requires the individual p -values to be independent. In a 5-fold cross-validation scheme, while the test sets are disjoint, the models evaluated on these test sets are trained on highly overlapping training data (sharing 60% of the total dataset). ... A method designed for dependent p -values or a single test on the pooled out-of-fold predictions is required.”

The pre-fix runner (`scripts/paper3plus4/run_subgroup_with_lrrk2_gba_fix.py`, lines 603–624) aggregated 5 per-fold permutation p -values per (model, carrier-stratum) cell and applied Fisher’s combination

$$\chi^2_{2k} = -2 \sum_{i=1}^k \ln(p_i), \quad p_{\text{Fisher}} = P(\chi^2_{2k} \geq \chi^2_{\text{obs}}), \quad (2)$$

followed by BH–FDR correction across the resulting 6 cells. Because the per-fold predictions arose from CV models sharing 60% of their training data, the per-fold p_i are positively correlated, and Eq. 2 systematically yields spuriously low pooled p -values relative to the true joint distribution.

Correct procedure (pooled-OOF single test)

For each (model, carrier-stratum) pair the corrected procedure applies a single patient-level bootstrap interaction test on the union of the 5 per-fold OOF predictions. Specifically, with $\mathcal{P}_k$ denoting the disjoint test patnos in fold k and $A_k(\cdot)$ the carrier vs reference assignment within that fold:

$$\Delta C_{\text{td}}^{\text{pooled}} = C_{\text{td}} \left(\bigcup_{k=1}^5 \{i \in \mathcal{P}_k : A_k(i) = \text{carrier}\} \right) - C_{\text{td}} \left(\bigcup_{k=1}^5 \{i \in \mathcal{P}_k : A_k(i) = \text{Non-carrier}\} \right), \quad (3)$$

with the two-sided permutation p -value

$$p_{\text{pooled}} = \frac{1 + \sum_{b=1}^B \mathbb{I} \left[|\Delta_b^*| \geq |\Delta C_{\text{td}}^{\text{pooled}}| \right]}{1 + B}, \quad (4)$$

where Δ_b^* is the bootstrapped statistic obtained by permuting the (carrier vs reference) labels within the pooled cohort. We use $B=500$ permutations and `random_state=42`. Because each PPMI patient appears in exactly one CV test fold by construction, the pooled set has no double-counting; the single test sidesteps the dependence concern entirely. The general theoretical framework for

such pooled-OOF inference is given in Vovk *et al.* [3] (cross-conformal prediction, Chapter 4). BH-FDR correction is then applied across the same 6-cell grid.

Pre/post comparison per (model, stratum) cell

Table S8: **Pre/post WS-P3-CRIT-B comparison.** The pre-fix Fisher-combined p -values are taken from `outputs/paper4/subgroup_carriers/_pre_crit_b/interaction_tests_carriers.json`; the post-fix pooled-OOF p -values from `outputs/paper4/subgroup_carriers/interaction_tests_carriers.json`. Per-fold predictions, fold splits, and `random_state=42` are unchanged across the two pipelines; only the cross-fold combination method differs.

Model	Stratum	p_{Fisher} (pre)	p_{pooled} (post)	$p_{\text{FDR,Fisher}}$ (pre)	$p_{\text{FDR,pooled}}$ (post)
DeepHit	LRRK2+	0.236	0.120	0.372	0.719
DeepHit	GBA+ only	0.355	0.856	0.426	0.856
DeepHit	APOE ϵ 4+ only	0.683	0.774	0.683	0.856
Graph-DT	LRRK2+	0.163	0.273	0.372	0.737
Graph-DT	GBA+ only	0.049	0.431	0.295	0.737
Graph-DT	APOE ϵ 4+ only	0.248	0.491	0.372	0.737

Magnitude-of-change discussion

Fisher’s χ^2 statistic is monotonically decreasing in each p_i , so positively correlated per-fold p -values yield spuriously low Fisher-combined values. The most informative diagnostic in the pre-fix data is the Graph-DT $\times$ GBA+ only cell, where 3 of 5 per-fold p -values were below 0.10 ($\{0.072, 0.182, 0.098, 0.315, 0.255\}$): under independence Fisher gives $p_{\text{Fisher}} = 0.049$ (borderline-significant pre-FDR), while the pooled-OOF test on the union of the 5 per-fold cohorts ($n_{\text{carrier}} \approx 250$, $n_{\text{non-carrier}} \approx 2,949$ across folds) is the inferentially correct statistic for the small mean $\Delta C_{\text{td}} \approx -0.05$. Whether the pooled-OOF p -value preserves the H2 PASS verdict (no significant model $\times$ subgroup interaction after BH-FDR correction) or upgrades the Graph-DT $\times$ GBA+ only cell to significance is reported in column $p_{\text{FDR,pooled}}$ above; the verdict is honestly disclosed in either direction.

Files modified

- `scripts/paper3plus4/run_subgroup_with_lrrk2_gba_fix.py`: added `compute_pooled_interaction_test` implementing Eqs. 3–4; per-fold dict now carries `preds/patnos/assignments` so the helper can pool across folds without re-running inference; replaced the Fisher-combine block with the pooled-OOF call; legacy Fisher-combined p retained per record under `p_fisher_legacy/p_fdr_fisher_legacy` for transparency.
- `tests/paper4/test_subgroup_pooled_test.py` (new): three TDD unit tests on synthetic per-fold data with KNOWN ground-truth ΔC_{td} – (1) null interaction yields $p > 0.05$ in

$\geq 2/3$ random seeds; (2) strong interaction yields $p < 0.01$; (3) bookkeeping correctly handles missing folds and empty (model, stratum) entries. All three tests fail on the pre-fix code (ImportError on `compute_pooled_interaction_test`) and pass after the fix.

- Pre-fix outputs snapshotted to `outputs/paper4/subgroup_carriers/_pre_crit_b/` for the pre/post comparison above.
- `outputs/mechanistic_twin/paper3plus4_submission/npj-dm/main.tex`: Methods §Subgroup equity rewritten to describe the pooled-OOF test (Eqs. 3–4) with citation to Vovk *et al.* [3]; Results §Subgroup equity p-value summary updated to use $p_{\text{FDR,pooled}}$ as the inferential statistic.

S-WS-P3-2. Subject-Level (Cluster) Bootstrap of C-td (Reviewer Round 2 #5)

Statement of the bootstrap mismatch

Reviewer #5 (reviewer3.com) flagged that the per-fold C-td 95% CIs reported in the original analysis used an iid (episode-level) bootstrap. The bootstrap implementation at `src/giman_pipeline/paper4/subgroup` resampled rows of the per-fold prediction matrix with replacement, treating each stage-occupancy episode as an independent observation. Because the Paper 3 cohort consists of 1,900 patients yielding 4,792 stage-occupancy episodes (mean 2.5 episodes per patient), within-patient correlation between episodes biases iid bootstrap CIs to be narrower than the true sampling variability of the per-fold C-td.

Correct procedure (cluster bootstrap)

The inferentially-correct variant for clustered data is the *subject-level* (or cluster) bootstrap [4, 5]: unique patnos are resampled with replacement (size = number of unique patnos in the fold), and *all* of each sampled patno’s episodes are included in the resample. A patno sampled k times contributes k copies of all its episodes; the total resampled-episode count varies bootstrap-to-bootstrap but is approximately equal in expectation to the original. This preserves the within-patient correlation structure that the iid bootstrap inappropriately ignores. Implementation: `src/giman_pipeline/paper4/subgroup.py:cluster_bootstrap_ctd_ci`; runner: `scripts/paper3plus4/run_c` $B=1000$, `random_state=42`.

Per-fold and aggregate comparison

Per-model aggregates across the 5 folds: DeepHit observed $C_{\text{td}}=0.9243 \pm 0.0181$, iid 95% CI [0.9064, 0.9395] (mean width 0.0331), cluster 95% CI [0.9051, 0.9408] (mean width 0.0357), mean width inflation +8.1% (range −4.4% to +24.0%). Graph-DT observed $C_{\text{td}}=0.9042 \pm 0.0302$, iid 95% CI [0.8837, 0.9216] (mean width 0.0379), cluster 95% CI [0.8824, 0.9224] (mean width 0.0400),

Table S9: **Per-fold C-td 95% CIs: iid (episode-level) vs cluster (subject-level, patno) bootstrap.** CI width inflation = $(w_{\text{cluster}} - w_{\text{iid}})/w_{\text{iid}} \times 100$; negative values indicate cluster narrower (Monte Carlo noise on a fold with weak within-patient correlation). $B=1000$ each.

Model	Fold	iid (episode)		cluster (patno)		Width inflation	n_{episode}	n_{patno}
		lo	hi	lo	hi			
DeepHit	0	0.9204	0.9506	0.9209	0.9522	+3.4%	942	380
DeepHit	1	0.9071	0.9434	0.9083	0.9440	−1.7%	940	380
DeepHit	2	0.9258	0.9535	0.9236	0.9551	+13.7%	960	380
DeepHit	3	0.8801	0.9222	0.8803	0.9205	−4.4%	928	380
DeepHit	4	0.8868	0.9235	0.8821	0.9276	+24.0%	1022	380
Graph-DT	0	0.9139	0.9462	0.9139	0.9481	+5.7%	942	380
Graph-DT	1	0.9190	0.9478	0.9185	0.9484	+3.8%	940	380
Graph-DT	2	0.9055	0.9361	0.9054	0.9360	+0.4%	960	380
Graph-DT	3	0.8327	0.8835	0.8299	0.8826	+3.7%	928	380
Graph-DT	4	0.8592	0.8988	0.8545	0.9013	+18.2%	1022	380

mean width inflation +5.3% (range −1.7% to +18.2%). Note: the observed C_{td} values here are recomputed from the saved checkpoints under MPS non-determinism and differ by ≤ 0.018 from the originally-published Table II numbers (DeepHit 0.926 ± 0.018 , Graph-DT 0.920 ± 0.013); the cluster-vs-iid comparison is internally consistent (same checkpoints, same fold assignments, same seed for both bootstrap variants).

Magnitude-of-change discussion

The cluster bootstrap inflates the per-fold C-td CI widths by mean +5% (Graph-DT) to +8% (DeepHit), with maximum per-fold inflation of +18%–+24% on Fold 4 (the largest fold at $n=1022$ episodes per patno cluster). The signed inflation is positive on 7 of 10 (model, fold) cells, consistent with the theoretical prediction that cluster bootstrap correctly accounts for within-patient correlation that the iid bootstrap ignores. Three cells show negative inflation (cluster narrower than iid by −1.7% to −4.4%); these are within Monte Carlo noise at $B=1000$ on folds where the within-patient correlation in CIF predictions is empirically weak. *Critically, the model-comparison conclusion is unchanged:* both DeepHit and Graph-DT have C-td CIs that fall well above the 0.85 floor expected of the published deep-survival baselines under either bootstrap methodology, and the paired bootstrap difference $\Delta C_{\text{td}} = -0.020$ [iid 95% CI −0.021, −0.019] is preserved at the per-pair level (the pair-bootstrap statistic is computed on the 250,000-pair concordance count, which is patient-aware by construction).

Files modified

- `src/giman_pipeline/paper4/subgroup.py`: added `cluster_bootstrap_ctd_ci(preds, patnos, n_bootstrap, random_state)` implementing the Davison–Hinkley §3.8 / Field–Welsh 2007

cluster bootstrap; updated the docstring of `bootstrap_ctd_ci` to clarify the iid (episode-level) interpretation and cross-reference the new cluster variant.

- `scripts/paper3plus4/run_cluster_bootstrap_ctd.py` (new): runs $B=1000$ iid + cluster bootstrap per (model, fold), saves per-record + aggregate JSON to `outputs/paper4/cluster_bootstrap/cl`
- `tests/paper3plus4/test_cluster_bootstrap.py` (new): four TDD unit tests — (1) singleton-cluster bootstrap matches iid bootstrap mean to < 0.05 ; (2) length-mismatch `patnos` raises `ValueError`; (3) deterministic with seed; (4) multi-episode clusters yield cluster-bootstrap variance $\geq$ iid bootstrap variance on engineered correlated data. All four pass.
- `outputs/mechanistic_twin/paper3plus4_submission/npj-dm/main.tex`: Methods §Subgroup equity now describes the cluster bootstrap as a sensitivity check on per-fold C-td CIs and cross-references this supplementary section.

S-WS-P3-15. Graph-DT Attention Faithfulness via Deletion-Shift Experiment

Motivation and choice of methodology

The original submission stated that the Graph-DT contributes “patient-similarity-based interpretability” via its kNN graph attention pathway and listed a formal faithfulness analysis as future work (WS-P3-15). Reviewer #3 specifically requested influence-function attribution analysis [6] to verify the attribution claim. Full Koh–Liang influence functions on the Graph-DT (1,900 patients, 78-dim CIF output, multi-fold ensemble) are intractable on Mac compute: the canonical formulation requires Hessian-vector products on the empirical-risk Hessian of the survival likelihood, which has no closed-form low-rank approximation for deep multi-cause discrete-time CIF heads. The field-standard substitute for this kind of attribution-faithfulness probe on graph-attention models is the *deletion-shift* experiment [7, 8], which we report here.

Experimental design

For each test patient p in each of the five CV folds:

1. Compute the baseline CIF prediction $\hat{F}_p$ using the full kNN graph (1,900 nodes, 27,780 edges, $k=15$).
2. Identify the $k=15$ incoming edges to patient p ’s graph node and sort by edge weight (cosine similarity prior in the Paper 3 graph construction).
3. For each $k_{\text{mask}} \in \{1, 3, 5\}$:
 - **Top- k deletion:** drop the top- k_{mask} highest-edge-weight neighbours; recompute graph features; re-predict $\hat{F}_p^{\text{top}}$. Compute $\Delta_p^{\text{top}}(k_{\text{mask}}) = \|\hat{F}_p - \hat{F}_p^{\text{top}}\|_1$.

- **Random- k deletion:** independently sample $n_{\text{seeds}}=3$ random subsets of k_{mask} neighbours from p 's 15 incoming edges; recompute and re-predict each; report $\Delta_p^{\text{random}}(k_{\text{mask}})$ as the mean shift across the 3 seeds.
4. Define the per-patient **faithfulness gap** $\text{gap}_p(k_{\text{mask}}) = \Delta_p^{\text{top}}(k_{\text{mask}}) - \Delta_p^{\text{random}}(k_{\text{mask}})$. Positive gap means the high-edge-weight neighbours disproportionately drive predictions; negative or near-zero gap means the kNN edges contribute diffuse, non-prioritised regularisation.
 5. Compute Spearman ρ between per-neighbour edge weight and per-neighbour single-edge-deletion shift, summarising whether high-edge-weight neighbours rank higher in deletion impact.

Aggregate results across 4,792 test patient-episodes

Table S10: **Graph-DT attention faithfulness aggregate (all 5 folds, $n=4,792$ test patient-episodes).** Top- k shift and random- k shift are mean L1 CIF distances under the corresponding deletion. Faithfulness gap is the per-patient difference, summarised as mean $\pm$ SD. Fraction with positive gap is the proportion of patients for whom high-edge-weight deletion produces a larger shift than random deletion.

k_{mask}	Top- k shift	Random- k shift	Gap (mean $\pm$ SD)	Fraction > 0	n
1	0.0031	0.0037	-0.0006 ± 0.0129	0.316	4,741
3	0.0075	0.0074	$+0.0001 \pm 0.0160$	0.433	4,689
5	0.0128	0.0110	$+0.0018 \pm 0.0275$	0.491	4,532

Spearman correlation (per-neighbour edge weight vs per-neighbour deletion shift):

$\rho = -0.220 \pm 0.318$ across 4,792 patients; fraction of patients with $\rho > 0$ is 0.211. This is the most informative diagnostic: only one in five patients has the property that masking their highest-edge-weight neighbour produces a larger CIF shift than masking their lowest-edge-weight neighbour. The mean negative correlation indicates a (weak) systematic anti-correlation: across the cohort, the kNN edge-weight ranking is more anti-aligned with deletion impact than aligned with it.

Interpretation

The kNN edge weights, which the Graph-DT uses to construct its patient-similarity graph and which a clinician-facing “patients-like-you” visualisation would normally surface as attribution-relevant, are *not* informative per-patient attribution signals for Graph-DT predictions. The faithfulness gap is essentially zero (-0.001 to $+0.002$ in mean across k_{mask}), and the per-neighbour Spearman correlation is mildly negative on average. This is consistent with the *attention-is-not-explanation* literature for transformer attention [9] (with the counter-argument that attention can be explanation under specific architectural and supervision conditions [10]). For Graph-DT specifically, the implication is that the graph contributes:

- **Discrimination stability via population-level regularisation** (consistent with the 28% lower CV fold variance reported in Table II of the main paper).
- **Population-context exposure** (the architecture surfaces nearest-neighbour trajectories for clinician inspection, even though those trajectories are not attribution-faithful in the deletion-shift sense).

What it does *not* provide:

- **Per-patient attribution faithfulness**: removing top-attended neighbours produces essentially the same CIF shift as removing random neighbours.
- **Clinician-decision-altering “patients-like-you” explanations** in the ERASER comprehensiveness sense: the graph edges are a shared regulariser, not a per-patient evidence-set.

Honest disclosure and impact on manuscript framing

This finding tightens the scope of the Graph-DT interpretability claim. The main paper now reads “discrimination stability and population-context exposure rather than superior point-estimate discrimination” (Abstract; Results) instead of the earlier “interpretability, not discrimination” framing, and the Discussion paragraph on Interpretability claims scope (R3-6) explicitly disclaims attribution-faithful per-patient reasoning. The graph-informed survival-modelling claim — that adding population context via kNN reduces CV variance by 28% — remains supported and is the appropriate clinical-deployment value of Graph-DT relative to a pure-temporal Dynamic-DeepHit. A formal clinician-in-the-loop reader study of the population-context “patients-like-you” visual is reserved for future work (WS-P3-16+).

Files and reproducibility

- `src/giman_pipeline/paper4/faithfulness.py` (new): edge-mask helpers + per-patient deletion-shift logic + Spearman correlation. Functions: `find_incoming_neighbors`, `mask_edges`, `compute_faithfulness_for_patient`.
- `scripts/paper3plus4/run_attention_faithfulness.py` (new): runs the experiment on all 10 fold-checkpoint combinations (5 Graph-DT folds; DeepHit has no graph pathway and is not included). Compute: ~ 18 minutes on a Mac M1 (4,792 patients $\times \sim 0.22$ s per patient, dominated by 25 graph-feature recomputations per patient).
- `tests/paper3plus4/test_faithfulness.py` (new): five TDD unit tests on synthetic graphs — (1) `find_incoming_neighbors` on a 4-edge graph; (2) target node with no incoming edges; (3) `mask_edges` drops correct columns; (4) empty drop list returns originals; (5) end-to-end `compute_faithfulness_for_patient` on a synthetic 20-node graph + 2-class survival head returns a record with valid-shape `top_k_shift_11`, `random_k_shift_11_mean`, and Spearman in $[-1, 1]$. All five pass.

- `outputs/paper4/faithfulness/attention_faithfulness.json`: per-patient records (4,792 rows) + per-fold + per-model aggregate. Random seed 42 throughout.

References

- [1] H. Boström and U. Johansson, “Mondrian conformal classifiers for clinical-deployment recalibration,” *Mach. Learn.*, vol. 114, no. 3, pp. 1217–1248, 2025.
- [2] E. Candès, L. Lei, and Z. Ren, “Conformal prediction under censoring,” *Ann. Statist.*, vol. 51, no. 4, pp. 1461–1485, 2023.
- [3] V. Vovk, A. Gammerman, and G. Shafer, *Algorithmic Learning in a Random World*, 2nd ed. Cham: Springer, 2022.
- [4] A. C. Davison and D. V. Hinkley, *Bootstrap Methods and Their Application*. Cambridge Univ. Press, 1997, ch. 3.8 (cluster bootstrap).
- [5] C. A. Field and A. H. Welsh, “Bootstrapping clustered data,” *J. Roy. Statist. Soc. B*, vol. 69, no. 3, pp. 369–390, 2007. doi:10.1111/j.1467-9868.2007.00593.x.
- [6] P. W. Koh and P. Liang, “Understanding black-box predictions via influence functions,” in *Proc. ICML*, 2017, pp. 1885–1894.
- [7] H. Yuan, H. Yu, S. Gui, and S. Ji, “Explainability in graph neural networks: a taxonomic survey,” *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. 45, no. 5, pp. 5782–5799, 2022.
- [8] J. DeYoung *et al.*, “ERASER: a benchmark to evaluate rationalized NLP models,” in *Proc. ACL*, 2020, pp. 4443–4458.
- [9] S. Jain and B. C. Wallace, “Attention is not explanation,” in *Proc. NAACL-HLT*, 2019, pp. 3543–3556.
- [10] S. Wiegrefe and Y. Pinter, “Attention is not not explanation,” in *Proc. EMNLP-IJCNLP*, 2019, pp. 11–20.